

EXECUTIVE MESSAGE

AAHP Swarm

From autonomous agents to accountable decisions

Review-first architecture for enterprise-grade AI orchestration.

Shift the conversation from AI that acts to AI whose decisions can be trusted.

AAHP Swarm operationalizes verification, not autonomy.

VALUE

Governance-ready decisions

Auditable AI decisions, reduced uncertainty, structured evidence.

DIFFERENTIATION

Evidence over narrative

Distributed roles, explicit checks, no implicit trust.

STRATEGIC FIT

Agentic DevOps foundation

EU AI Act readiness and scale across multi-agent stacks.

THE REVIEW CHAIN



SCOUT

Finds work, gaps and blockers



TESTER

Separates facts from assumptions



RISK

Evaluates risks and policy issues



VERDICT

Auditable decision: pass / warn / fail / block

Problem

- Self-validating agents create hidden risk
- Unclear what was actually checked
- Confidence based on narrative, not evidence
- No audit trail for AI-driven decisions

Insight

The bottleneck is verification, not generation. Quality comes from role separation and explicit evidence - not bigger models.

MARKET ANALYSIS

The verification gap in a \$50B+ opportunity

Agentic AI is exploding - but governance and review layers are the missing infrastructure.

AI AGENTS MARKET 2026

\$10.9B

Grand View Research, 2026

FORECAST 2030

\$50.3B

CAGR 45.8% (2025-2030)

AGENTIC AI 2034

\$139B

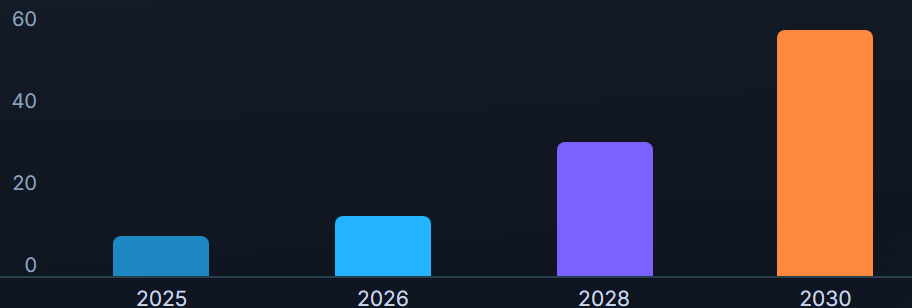
Fortune Business Insights

AI GOVERNANCE 2031

\$1.5B

CAGR 28.2% / EU AI Act

AI AGENTS MARKET SIZE 2025-2030 (USD BILLION)



Sources preserved from source deck: Grand View Research, Fortune Business Insights, McKinsey, Gartner, EU AI Act.

Key market drivers

- 62% of enterprises experimenting with AI agents; 23% scaling (McKinsey, 2025)
- 171% average ROI on deployed agents; top 5% return \$8 per \$1
- \$6.42B funding into agentic AI startups in 2025 alone
- 40%+ of agentic AI projects at risk of cancellation by 2027 due to poor scoping and governance
- EU AI Act enforcement from Feb 2025; fines up to EUR35M / 7% global turnover

Elvatis positioning

AAHP Swarm targets the verification and governance gap - turning failing pilots into auditable, deployable systems.

ROADMAP

From specification to autonomous operations

Review-first today. Governed automation tomorrow.

V0.1-V0.2

Foundations

Q2 2026

- Role contracts
- JSON schemas
- Public repo + license

V0.3

Review Runtime

Q3 2026

- Orchestration skeleton
- Deterministic verdict engine
- Audit-ready handoff

V0.4

Integrations

Q4 2026

- Claude / OpenAI / Gemini
- CI/CD + GitHub workflows
- Elvatis MCP integration

V1.0

Enterprise Ready

Q1 2027

- EU AI Act / ISO 42001 pack
- SSO, RBAC, audit logs
- Reference deployments

V2.0

Governed Autonomy

H2 2027

- Multi-swarm orchestration
- Self-improving review loops
- Verticals: Telco, Health, Public

Milestones & proof points

- Open-source community traction
- First enterprise pilot target: Q4 2026
- Compliance certification path target: Q2 2027
- Reference customer in EU public sector target: H2 2027

Strategic outcomes

- Trusted layer for agentic AI in regulated industries
- Foundation for Elvatis AI Governance suite
- Bridge to Deutsche Telekom / T-Systems / EU partners
- Defensible IP via contracts, schemas and verdict engine

COMPETITIVE ANALYSIS

Where AAHP Swarm plays

Orchestration frameworks deliver execution. AAHP Swarm delivers the review and governance layer they lack.

MARKET CONTEXT

Only 11% of enterprises run agentic AI in production; orchestration is consolidating, but governance remains open.

DIMENSION	LANGGRAPH	CREWAI	AUTOGEN / AG2	AAHP SWARM
Primary Focus	Stateful workflow graphs	Role-based agent crews	Conversational multi-agent	Review, verification and accountable handoff
Architecture	Directed state graphs	Declarative roles + tasks	Multi-turn message loops	Contract-driven swarm: Scout / Tester / Risk / Verdict
Production Strength	5/5 deterministic, audit-friendly	3/5 fast prototyping	4/5 MS / Azure ecosystem	5/5 review-first, audit-native
Governance Layer	Partial state, traces	Limited	Limited	Native - EU AI Act / ISO 42001 ready

AAHP Swarm is a layer

Sits above LangGraph, CrewAI and AutoGen; model-agnostic by design; adds review, verification and verdicts.

Strategic whitespace

Missing governance, EU AI Act demand, AI Governance CAGR 28.2%, first-mover in role-based review-first architecture.